

CDTM Drones Lecture 2016

Basics of Computer Vision

Part 3

CV, Object Detection, Tracking

Seyed-Ahmad Ahmadi

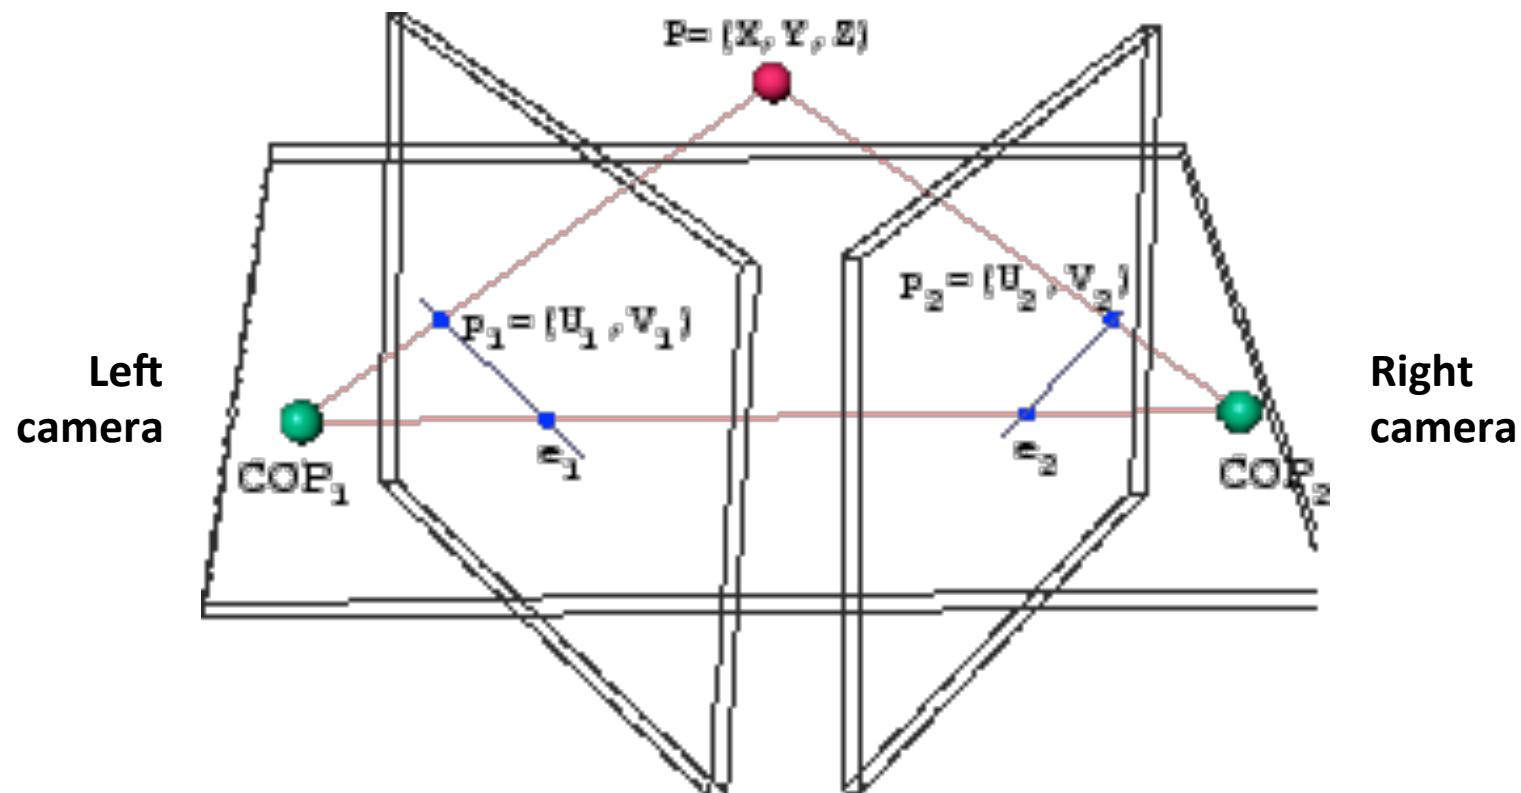
Dept. Neurology, LMU Klinikum Großhadern



Outline of this part

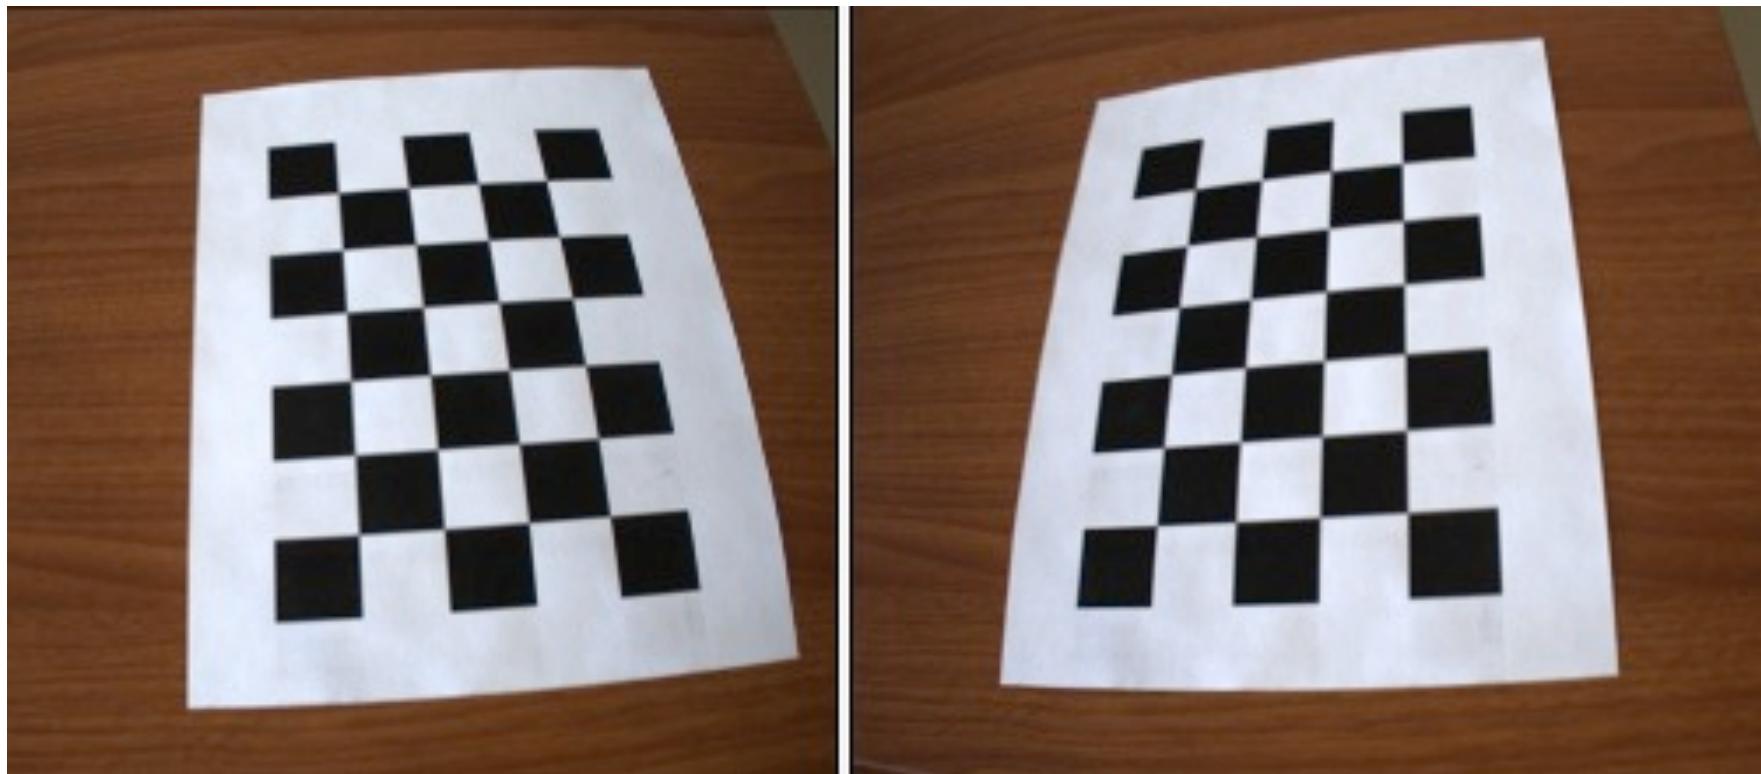
- Brief overview/history of object detection and tracking
- Image features for object tracking
- Basic concept of machine learning for object tracking

Historic motivation: Stereoscopic 3D vision



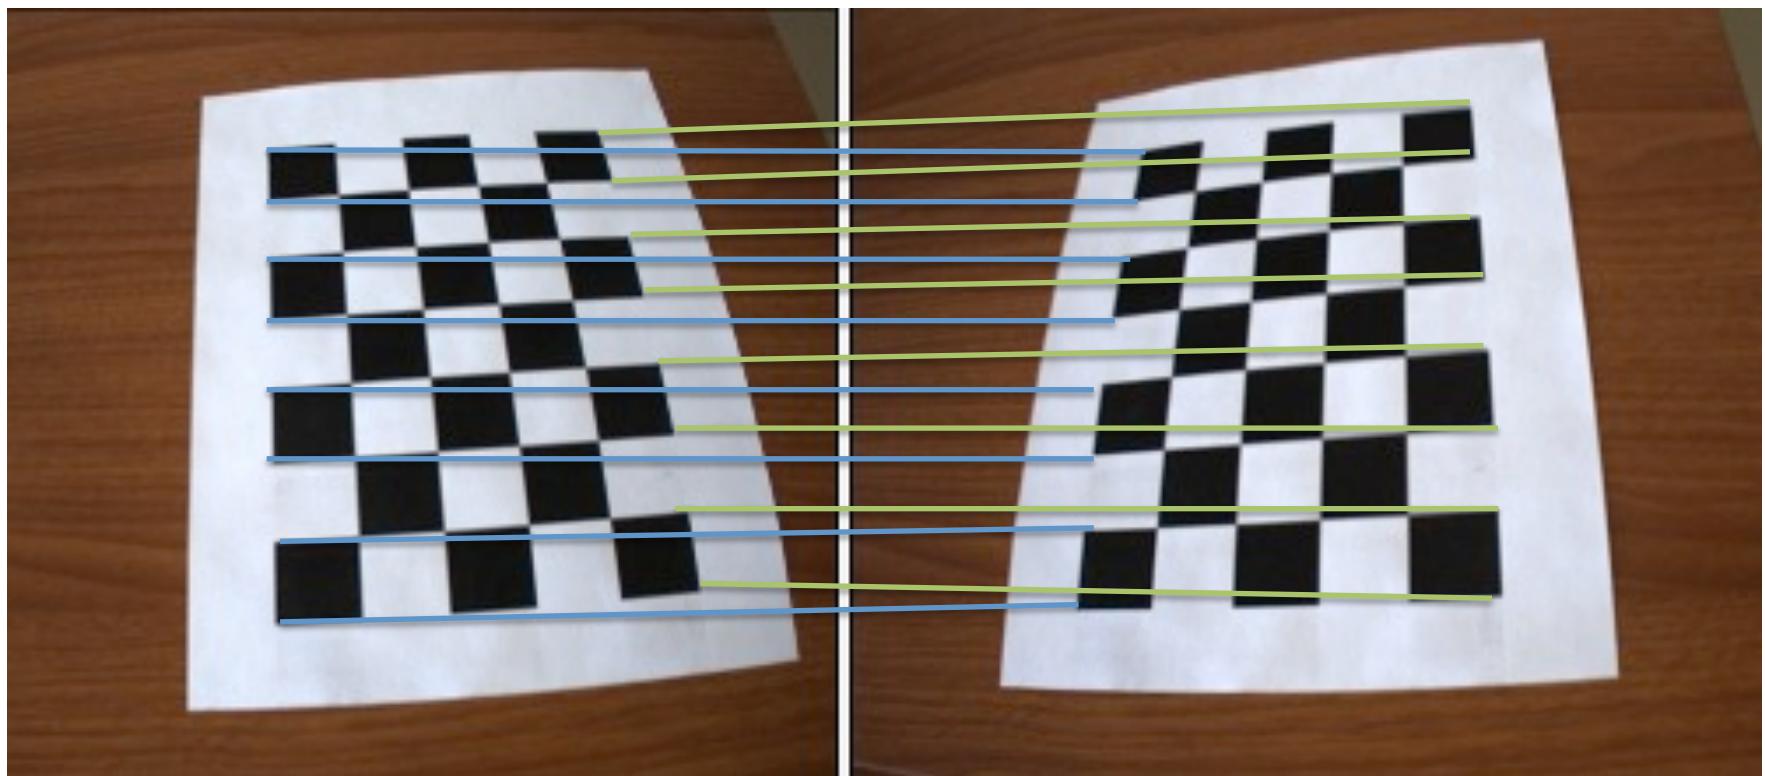
© National Instruments (NI), <http://www.ni.com/white-paper/14103/de/>

What the cameras see...



BoofCV real-time Java Computer Vision, http://boofcv.org/index.php?title=Example_Calibrate_Planar_Stereo

Establishing correspondence of key feature points (e.g. corners)

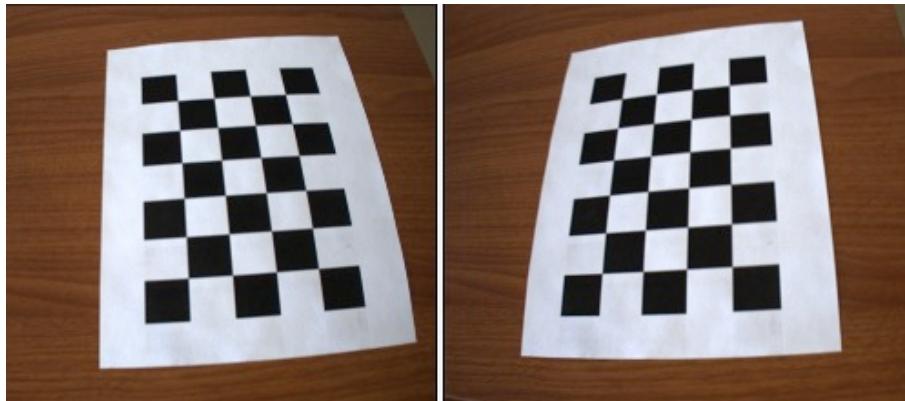


BoofCV real-time Java Computer Vision, http://boofcv.org/index.php?title=Example_Calibrate_Planar_Stereo

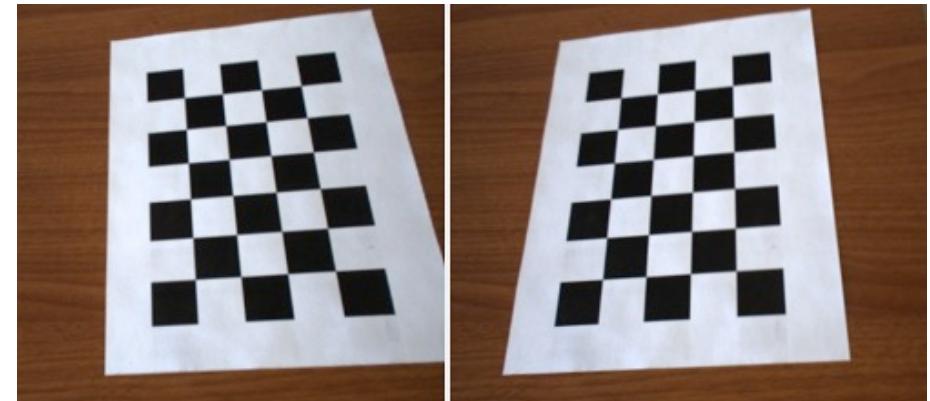
Using corresponding points, we can solve linear equations of camera parameters

- Intrinsic: focal length, image sensor format, and principal point
- Extrinsic: 3D rotation/translation from camera to world coordinates (position and heading)

Before calibration



After calibration



BoofCV real-time Java Computer Vision, http://boofcv.org/index.php?title=Example_Calibrate_Planar_Stereo

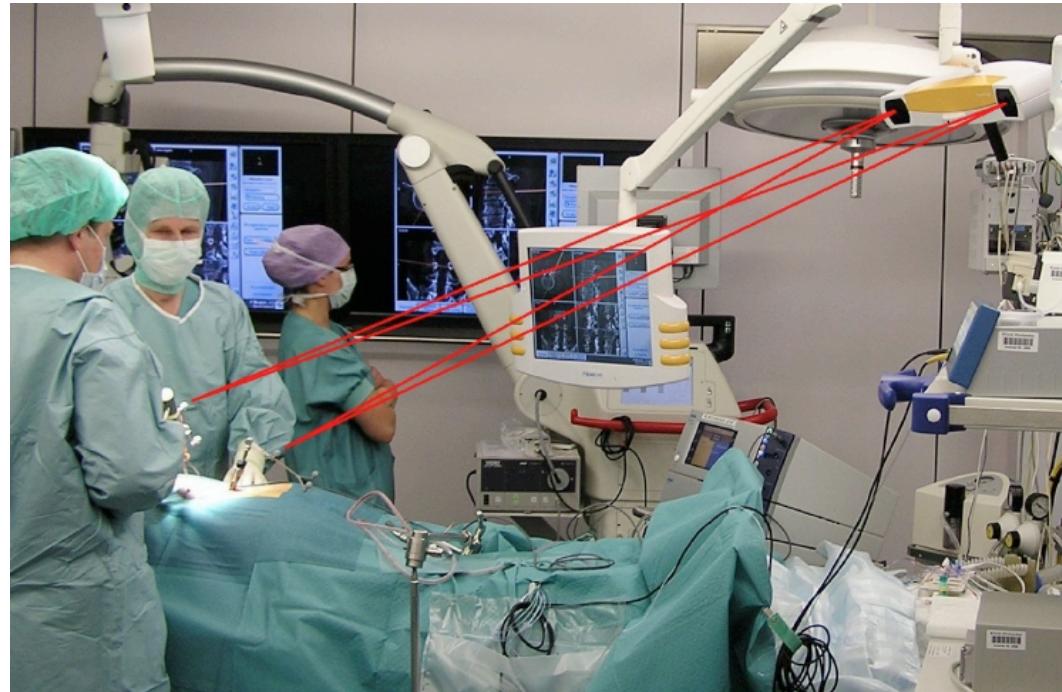
Applications of stereoscopic vision

Motion capturing
(Stereo- or multi-camera)



"MotionCapture", upload by T-tus
https://en.wikipedia.org/wiki/Motion_capture

Navigated surgery



<https://www.scienceopen.com/document/vid/03b46b7f-ebba-4229-b17d-b300d6ce5e7e?2>

Calibrated cameras allow for AR projections



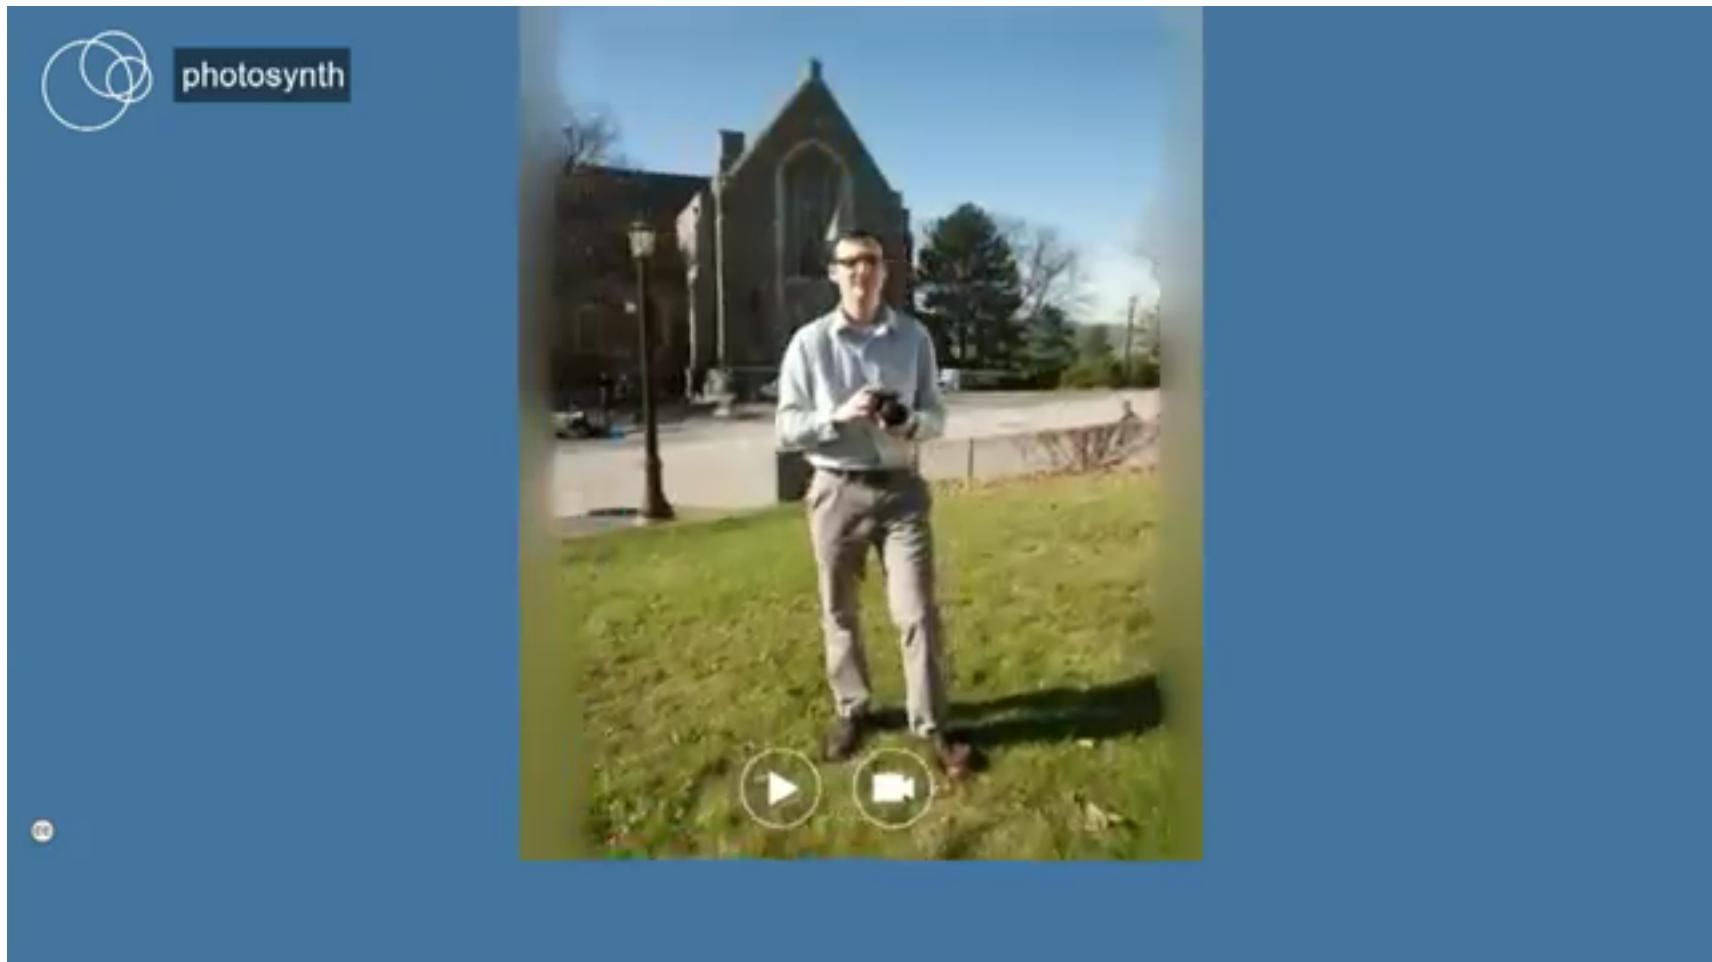
C. Bichlmeier, F. Wimmer, S.M. Heinig, N. Navab, Contextual Anatomic Mimesis: Hybrid In-Situ Visualization Method for Improving Multi-Sensory Depth Perception in Medical Augmented Reality, Int. Conf. ISMAR, 2007.

More than two cameras...



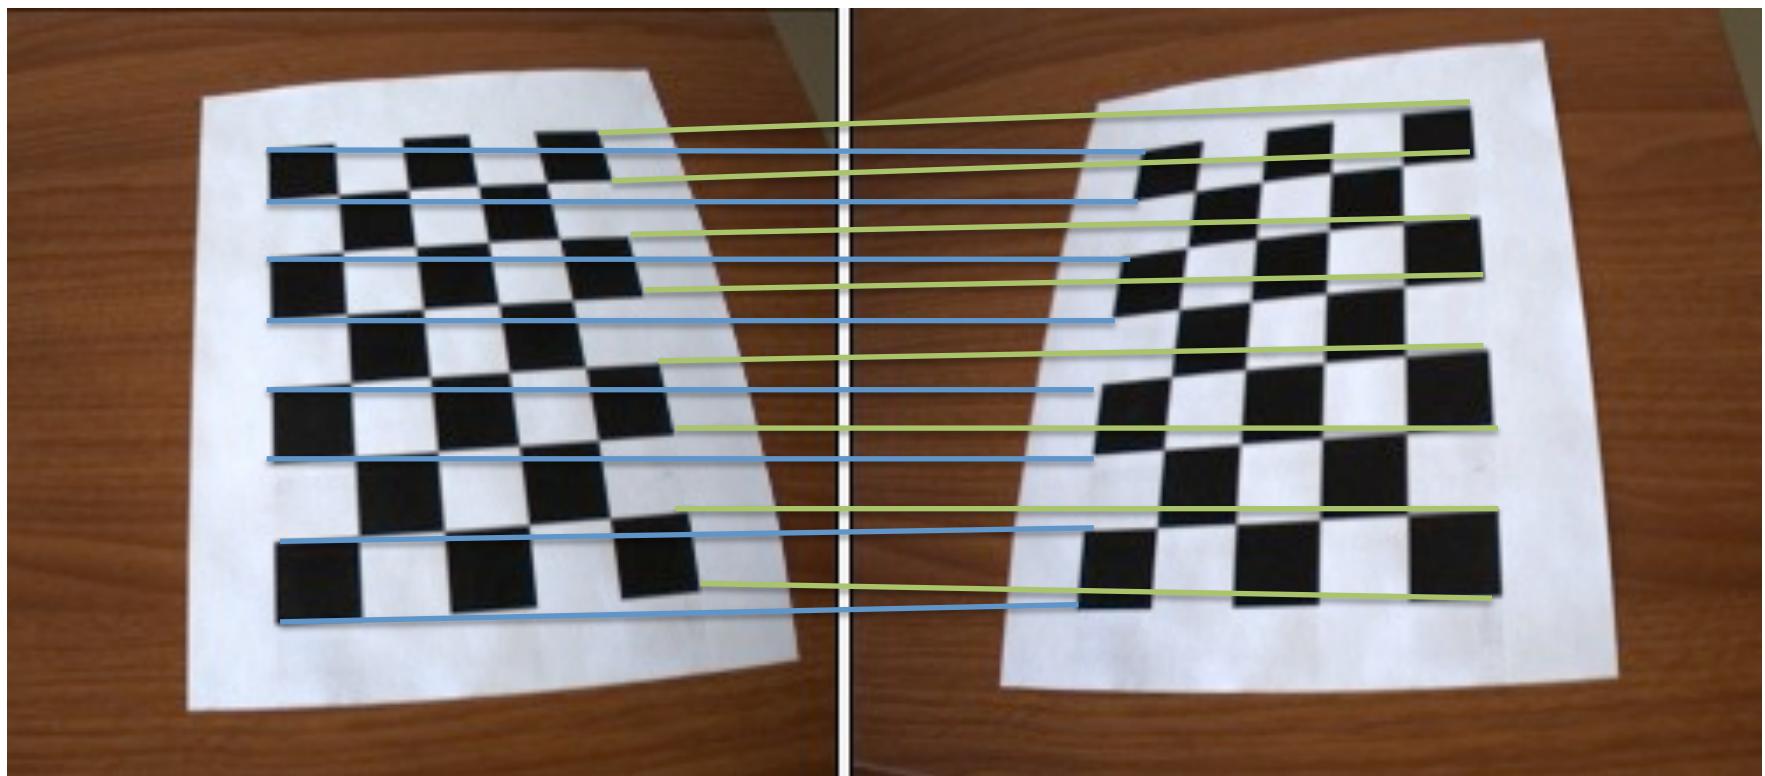
© Warner Bros., images from:
<http://sabia.tic.udc.es/gc/Contenidos%20adicionales/trabajos/Peliculas/FX/ej3.html>

Modern multi-view setups



©Microsoft Research, Photosynth Software, presented at TED, Dec. 2013
<https://www.youtube.com/watch?v=4LxlhoemR3A>

Establishing correspondence of key feature points (e.g. corners)



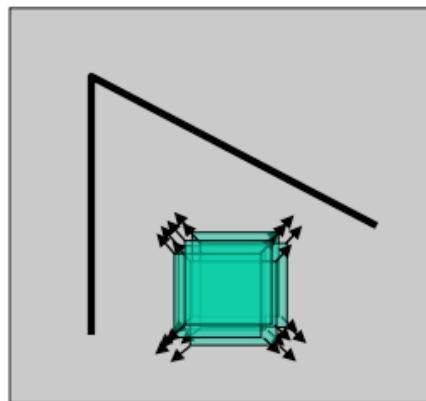
BoofCV real-time Java Computer Vision, http://boofcv.org/index.php?title=Example_Calibrate_Planar_Stereo

Features for tracking

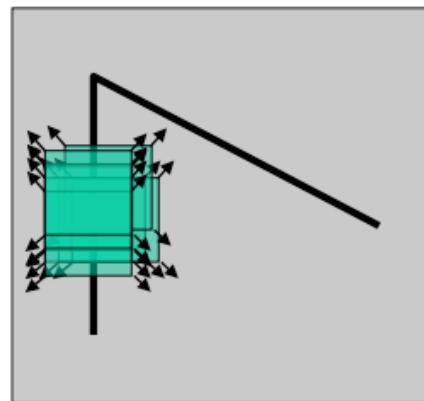
- Edges / Corners / Blobs...
 - Canny edge detector (1986, >22000 citations!)
 - Harris corner detector (1988, >11000 citations!)
 - FAST corner detector (2006, ~1800 citations)
- Shi-Tomasi
 - „Good features to track“ (1994, >6000 citations!)
- Feature detector/descriptors
 - Scale-Invariant-Feature-Transform (SIFT) (1999, ~31000 citations)
 - Speeded Up Robust Features (SURF) (2006, >6000 citations)

Harris Corner Detector

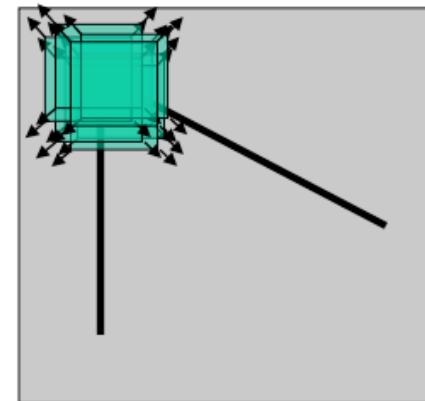
- Basic principle: corners are points of maximal change from one pixel to another



“flat” region:
no change in
all directions



“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions

<http://dsp.stackexchange.com/questions/14338/corner-detection-using-chris-harris-mike-stephens>

Harris Corner Detector

$$E(u, v) = \sum_{x,y} \underbrace{w(x, y)}_{\text{window function}} \frac{[\underbrace{I(x + u, y + v)}_{\text{shifted intensity}} - \underbrace{I(x, y)}_{\text{intensity}}]^2}{w(x, y)}$$

- Maximize difference between 2 imgs:
- Ends up in an eigenvalue analysis
- In OpenCV:
- `dst = cv2.cornerHarris(gray, 2, 3, 0.04)`

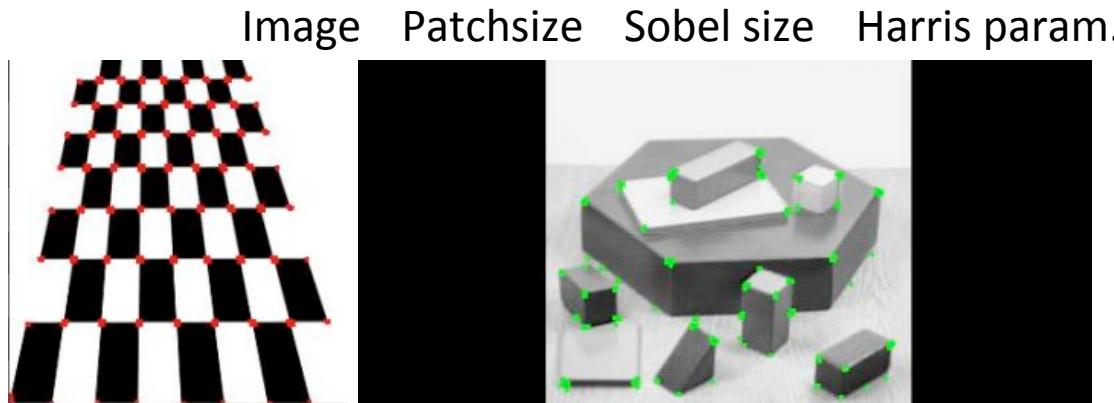
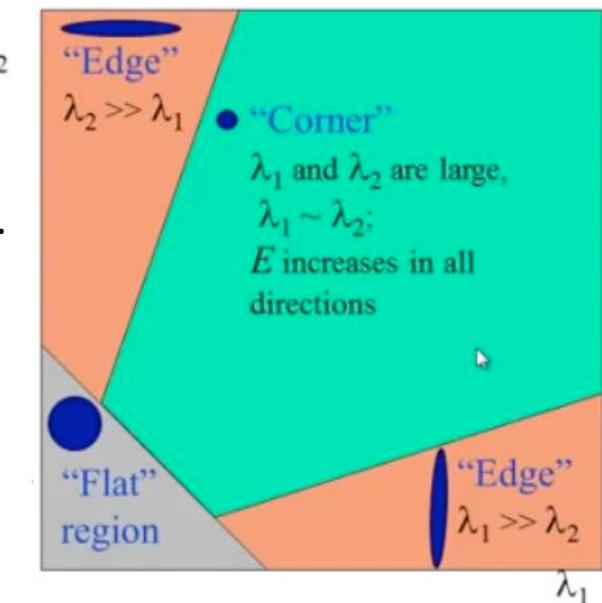


Image Patchsize Sobel size Harris param.

$$\frac{[\underbrace{I(x + u, y + v)}_{\text{shifted intensity}} - \underbrace{I(x, y)}_{\text{intensity}}]^2}{w(x, y)}$$

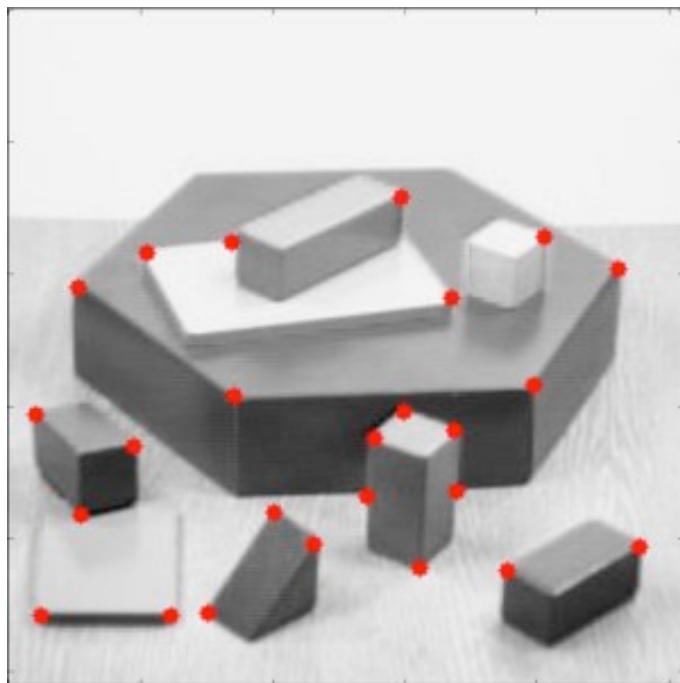


http://opencv-python-tutorials.readthedocs.org/en/latest/py_tutorials/py_feature2d/py_features_harris/py_features_harris.html

Shi Tomasi „Good Features to Track“

- In OpenCV:

```
cv2.goodFeaturesToTrack( ...  
    image, maxCorners, qualityLevel, minDistance)
```



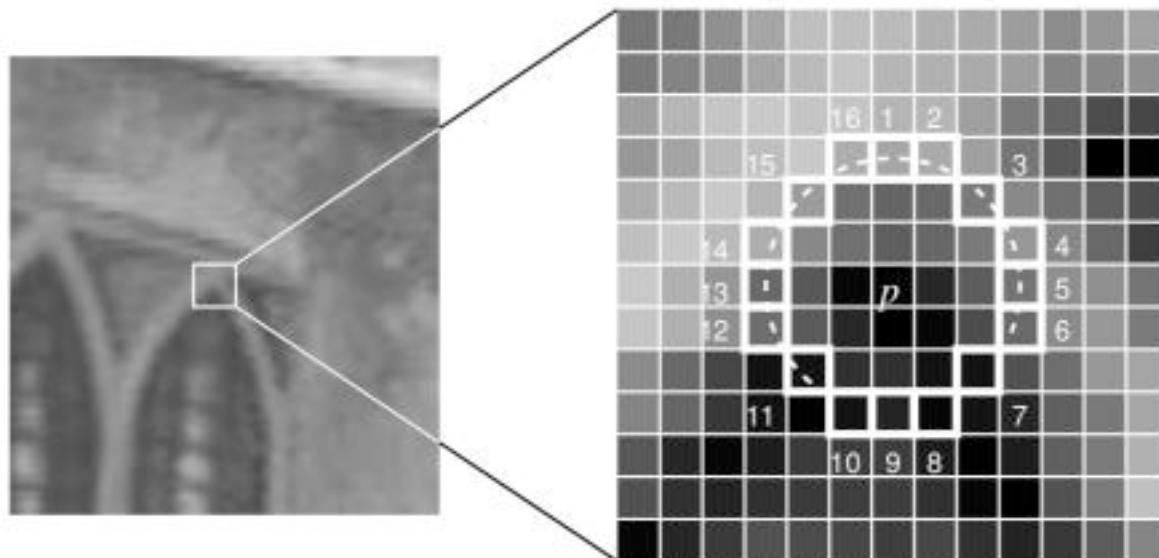
FAST Corner Detector

(Features from Accelerated Segment Test)

- Consider a circle around each pixel p
 - Corner = There must be at least n brighter/darker pixels than center pixel p
 - Even faster: only evaluate pixels at location 1-5-9-13
 - Weak/erroneous detections are compensated by domain-specific machine learning
 - Adjacent detections are compensated with non-maximum suppression
 - Several times faster than other corner detectors

```
kp = fast.detect(img,None)
img2 = cv2.drawKeypoints(img, kp, color=(255,0,0))
```

```
fast = cv2.FastFeatureDetector()
kp = fast.detect(img, None)
img2 = cv2.drawKeypoints(img, kp, color=(255, 0, 0))
```



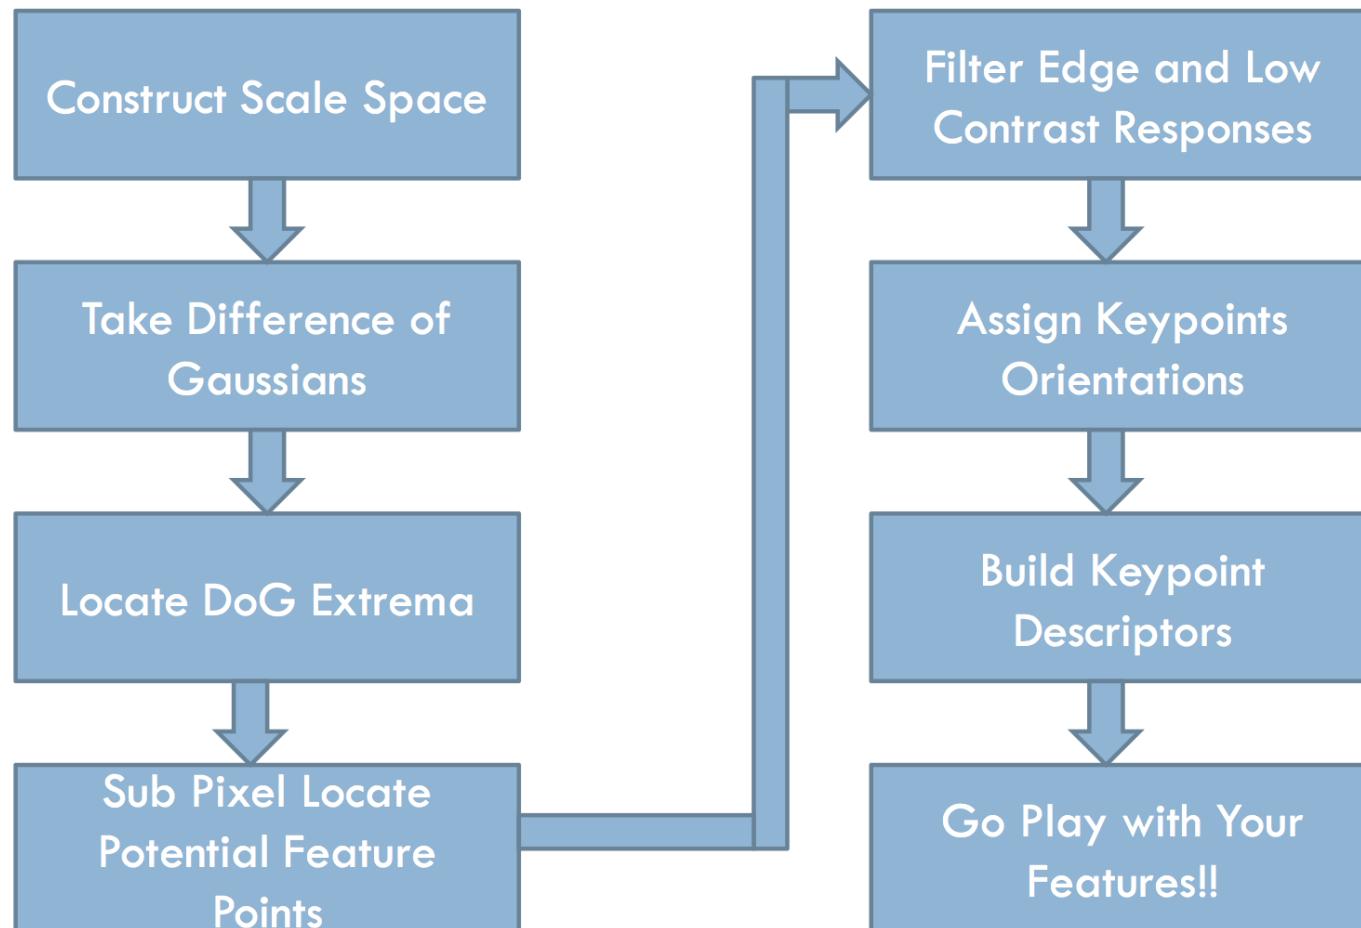
SIFT Feature Detector+Descriptor

(Scale Invariant Feature Transform)

>31.000 citations!!

- Invented in 1999 by David Lowe, patented!!
- Many advantages
 - Scale invariance
 - Rotation invariance
 - Illumination invariance
 - Viewpoint invariance (to some degree)

SIFT Feature Detector+Descriptor (Scale Invariant Feature Transform)



© Jason Clemons, U.Mich.

http://web.eecs.umich.edu/~silvio/teaching/EECS598/lectures/lecture10_1.pdf

SIFT in Python/OpenCV

```
# first, detect keypoints
sift = cv2.SIFT()
kp = sift.detect(gray,None)
img=cv2.drawKeypoints(gray,kp,flags= cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
cv2.imshow('features',img)

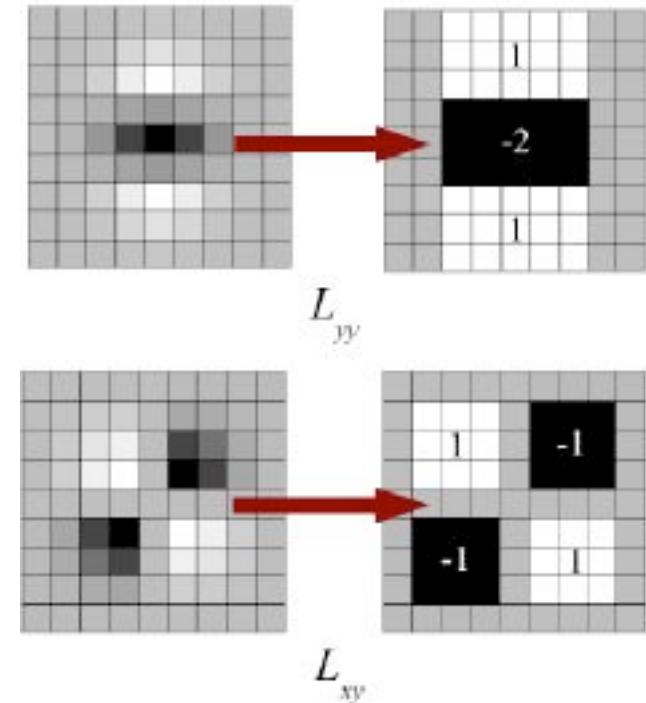
# second, compute the keypoint descriptors
kp,des = sift.compute(gray,kp)

# or, do it in one pass
kp, des = sift.detectAndCompute(gray,None)
```



SURF Feature Detector+Descriptor (Speeded Up Robust Features)

- ~3x faster than SIFT (more suited for real-time applications)
- Similarly accurate as SIFT
- Main reason for being faster:
 - Approximates Laplacian-of-Gaussian via box filters, rather than Difference-of-Gaussians in scale-space
 - Uses local wavelet filters responses as descriptors
- SURF is good at handling images with blurring and rotation, but not good at handling viewpoint change and illumination change.



SURF in Python/OpenCV

```
# create SURF object, with Hessian threshold 400
# (the higher the threshold, the fewer keypoints are found; in practice, use 300-500)
surf = cv2.SURF()

# Find keypoints and descriptors directly
kp, des = sift.detectAndCompute(gray,None)

# Draw keypoints on image
img2 = cv2.drawKeypoints(img,kp,None,(255,0,0),4)
plt.imshow(img2),plt.show()

# Optional: compute un-oriented SURF
# Even faster than SURF, but:
# Orientation invariance only up to +/- 15°
surf.upright = True
kp = surf.detect(img,None)
```



Many more descriptors and techniques

- Histogram of Oriented Gradients (HoG)
- Local Invariant Descriptors
- BRIEF
- ORB (Oriented FAST, Rotated BRIEF)
- Bag-of-Visual-Words
- Deformable Parts Models
- Exemplar Models
- Etc. etc...
- **30-40 years of computer vision research**

Matching of images via features

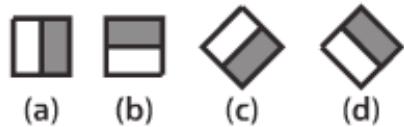
1. Find good features/keypoints in image A
2. Find good features/keypoints in image B
3. Try to match the two sets of keypoints with a transformation model
 - Robust matching despite noisy features and noisy matches
 - Specialized algorithms (e.g. Random Sample Consensus, RANSAC)

Object detection in images via machine learning

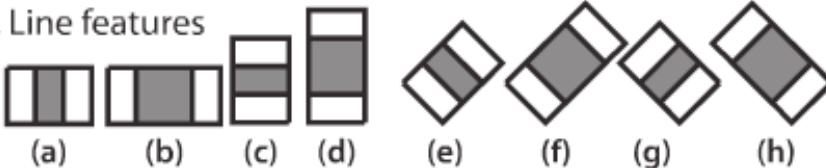
1. Image-wise object detection, no tracking
2. Extract good features at all image locations
3. Train a machine learning algorithm to find the object

Face Detection Using Haar Cascades

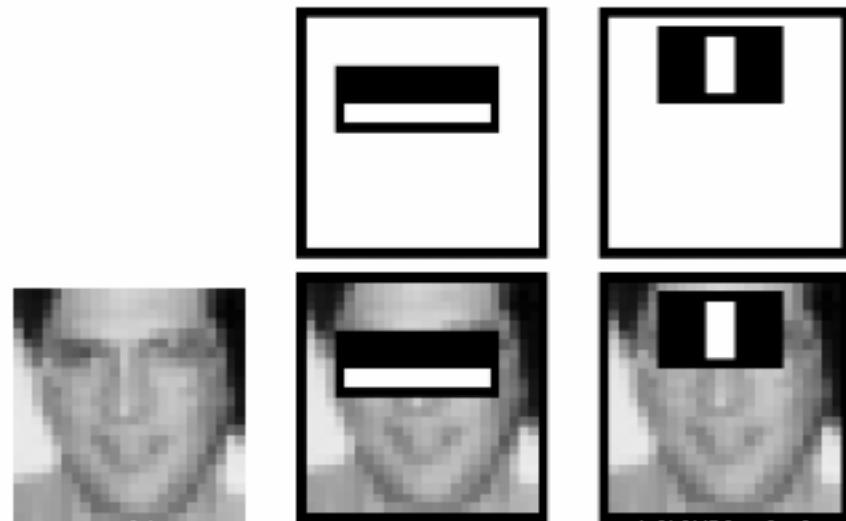
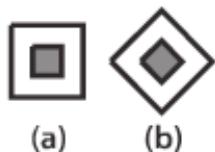
1. Edge features



2. Line features

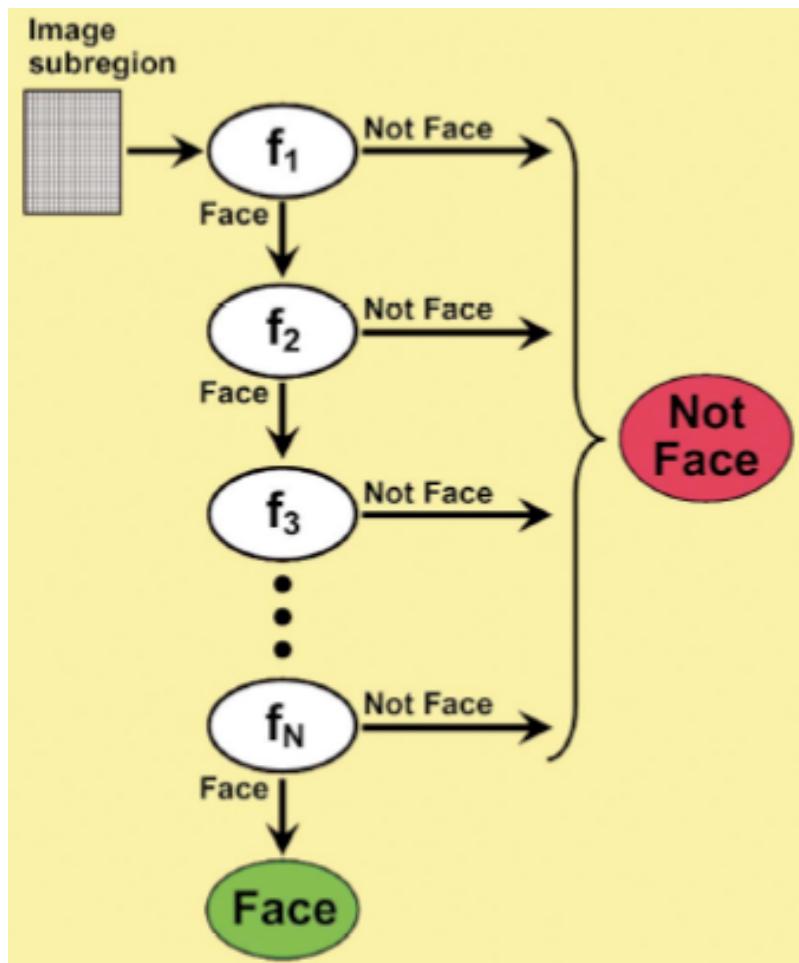


3. Center-surround features



- >160.000 features! How to choose best ones?
- Adaboost (Ensemble machine learning)

Cascade of Classification



```
haarcascade_eye_tree_eyeglasses.xml  
haarcascade_eye.xml  
haarcascade_frontalface_alt2.xml  
haarcascade_frontalface_alt_tree.xml  
haarcascade_frontalface_alt.xml  
haarcascade_frontalface_default.xml  
haarcascade_fullbody.xml  
haarcascade_lefteye_2splits.xml  
haarcascade_lowerbody.xml  
haarcascade_mcs_eyepair_big.xml  
haarcascade_mcs_eyepair_small.xml  
  
haarcascade_mcs_leftear.xml  
haarcascade_mcs_lefteye.xml  
haarcascade_mcs_mouth.xml  
haarcascade_mcs_nose.xml  
haarcascade_mcs_rightear.xml  
haarcascade_mcs_righteye.xml  
haarcascade_mcs_upperbody.xml  
haarcascade_profileface.xml  
haarcascade_righteye_2splits.xml  
haarcascade_smile.xml  
haarcascade_upperbody.xml
```

The Code

```
import numpy as np
import cv2
from matplotlib import pyplot as plt

face_cascade = cv2.CascadeClassifier('haarcascade_frontalface_default.xml')
eye_cascade = cv2.CascadeClassifier('haarcascade_eye.xml')

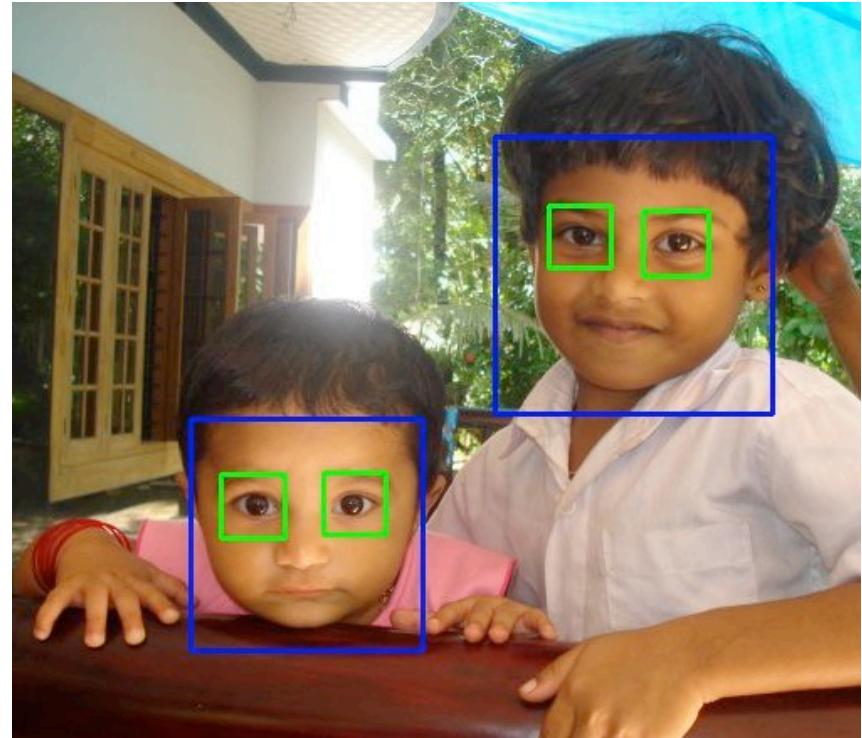
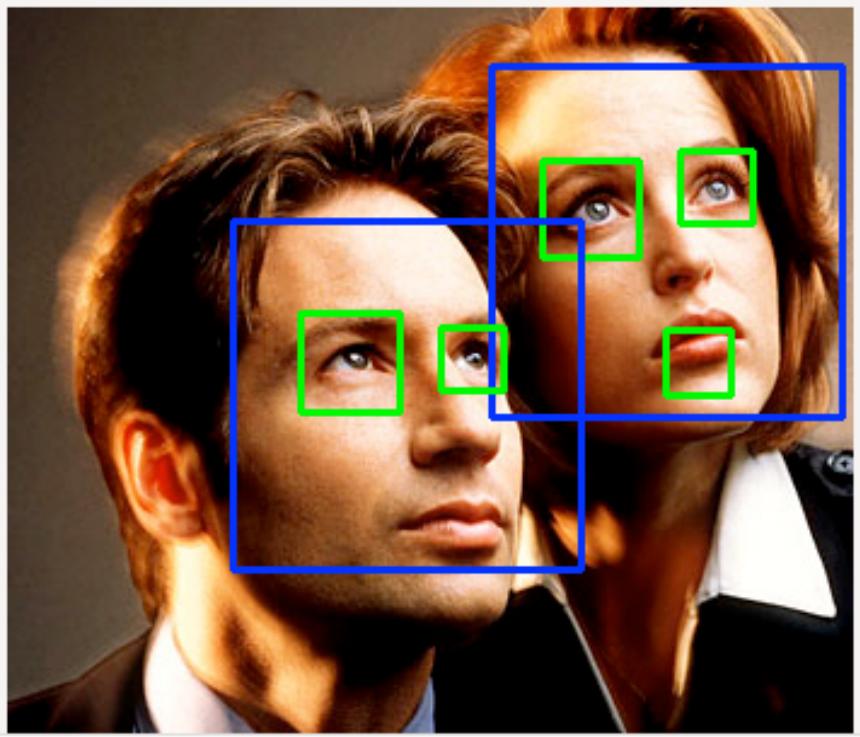
img = cv2.imread('xfiles4.jpg')
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

faces = face_cascade.detectMultiScale(gray, 1.3, 5)

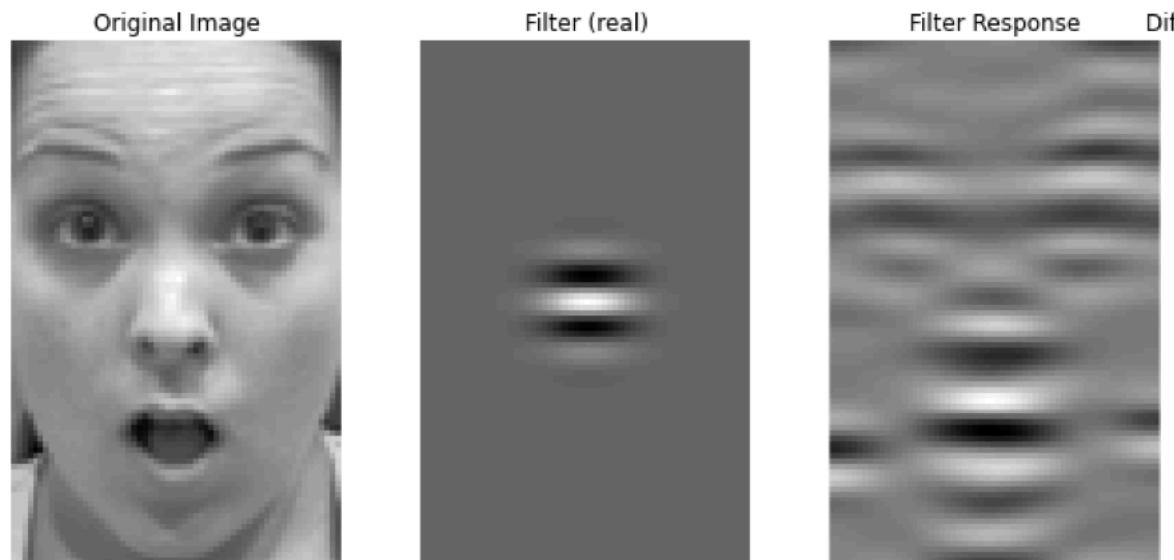
for (x,y,w,h) in faces:
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)
    roi_gray = gray[y:y+h, x:x+w]
    roi_color = img[y:y+h, x:x+w]
    eyes = eye_cascade.detectMultiScale(roi_gray)
    for (ex,ey,ew,eh) in eyes:
        cv2.rectangle(roi_color,(ex,ey),(ex+ew,ey+eh),(0,255,0),2)

cv2.imshow('img',img)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

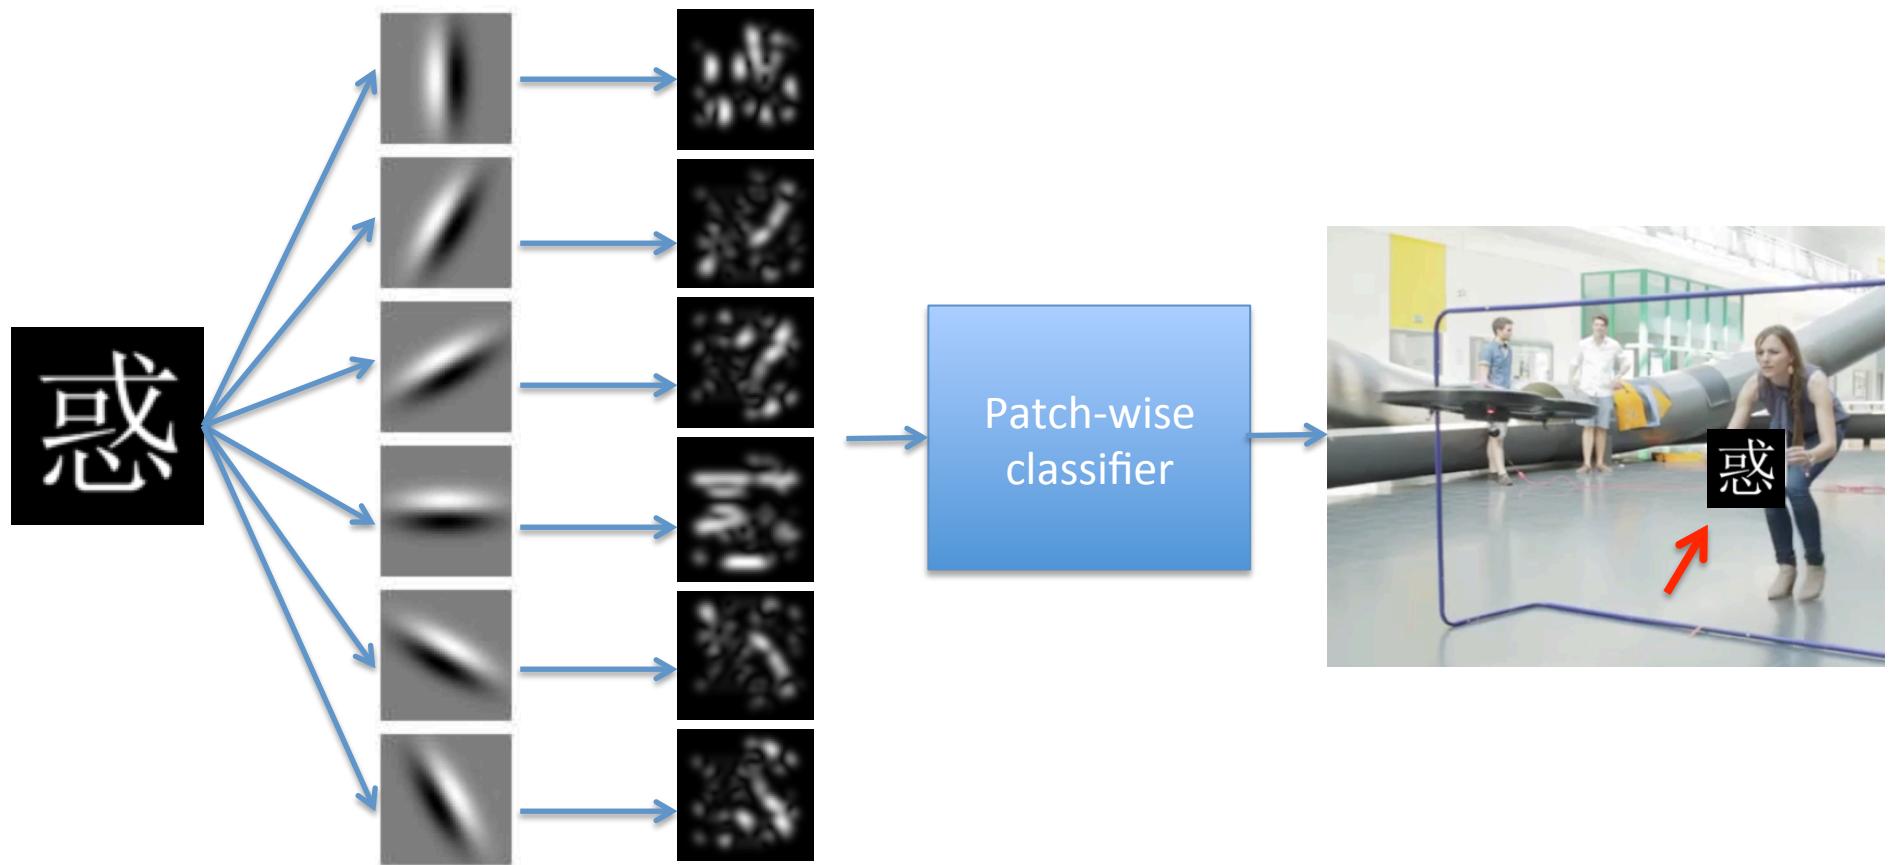
Results



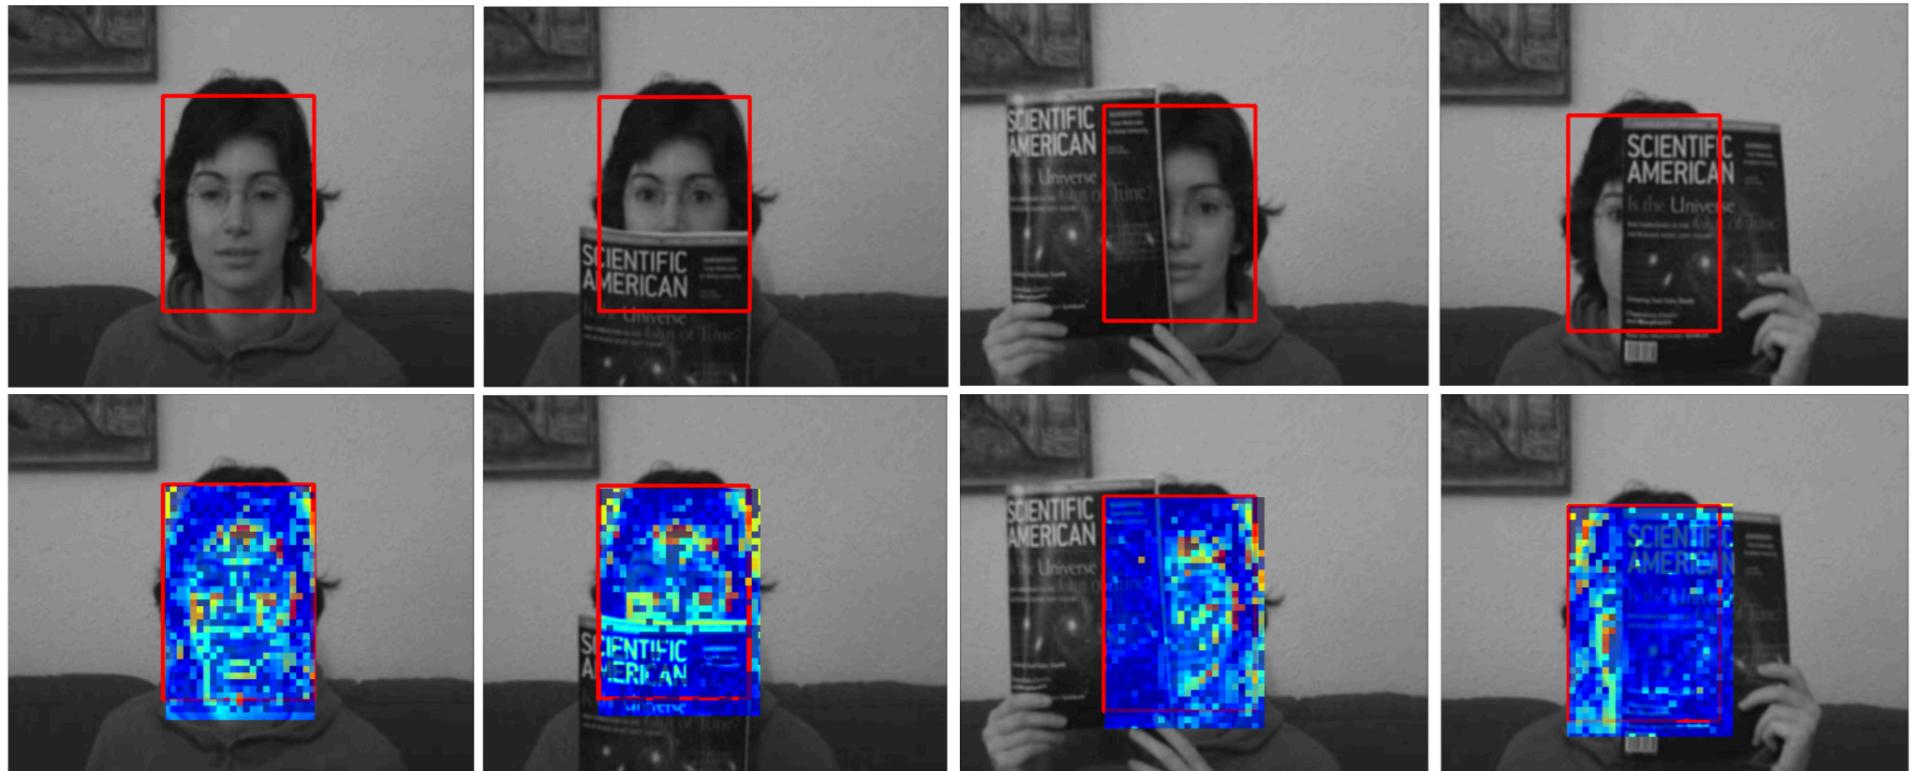
Filters can represent „good features“ in images



Gabor filter banks mimic visual cortex basic functions



Filter responses for local specificity and global robustness



Conclusion

- Computer vision was motivated by stereo 3D vision and matching of keypoints between image pairs
- Keypoints need to be identified via features
 - 30-40 years of computer vision community to find good keypoints and good matching algorithms
- Features can be extracted for each pixel via banks of filters
- Local feature responses from many filters give a good clue (e.g. to a classifier) where the object is

Questions?